

The full_join verb

JOINING DATA WITH DPLYR



Chris Cardillo

Data Scientist at DataCamp

Left and right joins


```
batwing %>%  
  left_join(batwing, by = c("part_num", "color_id"), suffix = c("_batmobile", "_batwing"))
```

```
# A tibble: 309 x 4  
  part_num color_id quantity_batmobile quantity_batwing  
  <chr>      <dbl>          <dbl>          <dbl>  
1 3023         0             22             22  
2 3024         0             22             22  
3 3623         0             20             20  
4 11477        0             18             18  
5 99207        71             18             18  
6 2780         0             17             17  
7 3666         0             16             16  
8 22385        0             14             14  
9 3710         0             14             14  
10 99563        0             13             13  
# ... with 299 more rows
```

The full join

first table

second table



Joining and filtering

```
inventory_parts_joined <- inventories %>%  
  inner_join(inventory_parts, by = c("id" = "inventory_id")) %>%  
  arrange(desc(quantity)) %>%  
  select(-id, -version)
```

```
batmobile <- inventory_parts_joined %>%  
  filter(set_num == "7784-1") %>%  
  select(-set_num)
```

```
batwing <- inventory_parts_joined %>%  
  filter(set_num == "70916-1") %>%  
  select(-set_num)
```

Batmobile vs. Batwing

batmobile

```
# A tibble: 173 x 3
  part_num color_id quantity
  <chr>      <dbl>     <dbl>
1 3023         72         62
2 2780          0         28
3 50950         0         28
4 3004         71         26
5 43093          1         25
6 3004          0         23
7 3010          0         21
8 30363         0         21
9 32123b        14         19
10 3622          0         18
# ... with 163 more rows
```

batwing

```
# A tibble: 309 x 3
  part_num color_id quantity
  <chr>      <dbl>     <dbl>
1 3023          0         22
2 3024          0         22
3 3623          0         20
4 11477         0         18
5 99207         71         18
6 2780          0         17
7 3666          0         16
8 22385         0         14
9 3710          0         14
10 99563         0         13
# ... with 299 more rows
```

Joining it all together

Left join: keep all batmobile

```
batmobile %>%  
  left_join(batwing, by = c("part_num", "color_id"), suffix = c("_batmobile", "_batwing"))
```

Right join: keep all batwing

```
batmobile %>%  
  right_join(batwing, by = c("part_num", "color_id"), suffix = c("_batmobile", "_batwing"))
```

Full join: keep all both

```
batmobile %>%  
  full_join(batwing, by = c("part_num", "color_id"), suffix = c("_batmobile", "_batwing"))
```

Full join result

```
batmobile %>%  
  full_join(batwing, by = c("part_num", "color_id"), suffix = c("_batmobile", "_batwing"))
```

```
# A tibble: 440 x 4  
  part_num color_id quantity_batmobile quantity_batwing  
  <chr>      <dbl>          <dbl>          <dbl>  
1 3023         72             62             NA  
2 2780          0             28             17  
3 50950         0             28              2  
4 3004         71             26              2  
5 43093          1             25              6  
6 3004          0             23              4  
7 3010          0             21             NA  
8 30363         0             21             NA  
9 32123b        14             19             NA  
10 3622          0             18              2  
# ... with 430 more rows
```

Replace NA: multiple variables

```
library(tidyr)

batmobile %>%
  full_join(batwing, by = c("part_num", "color_id"), suffix = c("_batmobile", "_batwing")) %>%
  replace_na(list(quantity_batmobile = 0,
                  quantity_batwing = 0))
```


Let's practice!

JOINING DATA WITH DPLYR

The semi- and anti-join verbs

JOINING DATA WITH DPLYR



Chris Cardillo
Data Scientist at DataCamp

Mutating verbs

- `inner_join`
- `left_join`
- `right_join`
- `full_join`

Review: left join

```
batmobile %>%  
  left_join(batwing, by = c("part_num", "color_id"), suffix = c("_batmobile", "_batwing"))
```

```
# A tibble: 173 x 4  
  part_num color_id quantity_batmobile quantity_batwing  
  <chr>      <dbl>          <dbl>          <dbl>  
1 3023      72            62            NA  
2 2780      0             28             17  
3 50950     0             28              2  
4 3004      71            26              2  
5 43093     1             25              6  
6 3004      0             23              4  
7 3010      0             21             NA  
8 30363     0             21             NA  
9 32123b    14             19             NA  
10 3622     0             18              2  
# ... with 163 more rows
```


Filtering joins

- Keeps or removes observations from the first table
- Doesn't add new variables
- `semi_join()`
- `anti_join()`

Filtering joins

Semi join


- What observations in X are **also** in Y?



Filtering joins


Semi join

- What observations in X are **also** in Y?



Anti join

- What observations in X are **not** in Y?



The semi join

```
batmobile %>%  
  semi_join(batwing, by = c("color_id", "part_num"))
```

```
# A tibble: 45 x 3  
  part_num color_id quantity  
  <chr>      <dbl>     <dbl>  
1 2780         0         28  
2 50950        0         28  
3 3004         71         26  
4 43093         1         25  
5 3004         0         23  
6 3622         0         18  
7 4286         0         16  
8 3039         0         12  
9 4274         71         12  
10 3001         0         11  
# ... with 35 more rows
```


The anti join

```
batmobile %>%  
  anti_join(batwing, by = c("color_id", "part_num"))
```

```
# A tibble: 128 x 3  
  part_num color_id quantity  
  <chr>      <dbl>     <dbl>  
1 3023         72         62  
2 3010          0         21  
3 30363         0         21  
4 32123b        14         19  
5 50950        320         18  
6 6541          0         18  
7 3040b         0         14  
8 3298          0         14  
9 3660          0         14  
10 42022         0         14  
# ... with 118 more rows
```

Filtering with semi join

```
themes %>%  
  semi_join(sets, by = c("id" = "theme_id"))
```

```
# A tibble: 569 x 3  
   id name      parent_id  
  <dbl> <chr>      <dbl>  
1     1 Technic          NA  
2     2 Arctic Technic      1  
3     3 Competition        1  
4     4 Expert Builder      1  
5     5 Model                1  
6     6 Airport              5  
7     7 Construction        5  
8     9 Fire                 5  
9    10 Harbor              5  
10    11 Off-Road            5  
# ... with 559 more rows
```

Filtering with anti join

```
themes %>%  
  anti_join(sets, by = c("id" = "theme_id"))
```

```
# A tibble: 96 x 3  
   id name      parent_id  
  <dbl> <chr>      <dbl>  
1     8 Farm         5  
2    24 Airport      23  
3    25 Castle      23  
4    26 Construction 23  
5    27 Race         23  
6    28 Harbor       23  
7    29 Train        23  
8    32 Robot        23  
9    34 Building     23  
10   35 Cargo        23  
# ... with 86 more rows
```

The joining verbs

- `inner_join`
- `left_join`
- `right_join`
- `full_join`
- `semi_join`
- `anti_join`

Let's practice!

JOINING DATA WITH DPLYR

Visualizing set differences

JOINING DATA WITH DPLYR



Chris Cardillo
Data Scientist at DataCamp

Aggregating sets into colors

```
batmobile_colors <- batmobile %>%  
  group_by(color_id) %>%  
  summarize(total = sum(quantity))
```

```
batmobile_colors
```

```
# A tibble: 12 x 2  
  color_id total  
    <dbl> <dbl>  
1         0   543  
2         1    33  
3         4    16  
4        14    20  
5        15    16  
6        36    15  
7        57     8  
8        71   202  
9        72   160  
10       182     8  
# ... with 2 more rows
```

```
batwing_colors <- batwing %>%  
  group_by(color_id) %>%  
  summarize(total = sum(quantity))
```

```
batwing_colors
```

```
# A tibble: 20 x 2  
  color_id total  
    <dbl> <dbl>  
1         0   418  
2         1    45  
3         4    81  
4        14    22  
5        15    22  
6        19    10  
7        25     1  
8        34     3  
9        36     9  
10       46    21  
# ... with 10 more rows
```

Comparing color schemes of sets

```
batmobile_colors %>%  
  full_join(batwing_colors, by = "color_id", suffix = c("_batmobile", "_batwing")) %>%  
  replace_na(list(total_batmobile = 0, total_batwing = 0))
```

```
# A tibble: 22 x 3  
  color_id total_batmobile total_batwing  
    <dbl>         <dbl>         <dbl>  
1         0           543           418  
2         1           33            45  
3         4           16            81  
4        14           20            22  
5        15           16            22  
6        36           15             9  
7        57            8             3  
8        71          202          158  
9        72          160          213  
10       182            8           14  
# ... with 12 more rows
```


Adding the color names

```
batmobile_colors %>%  
  full_join(batwing_colors, by = "color_id", suffix = c("_batmobile", "_batwing")) %>%  
  replace_na(list(total_batmobile = 0, total_batwing = 0)) %>%  
  inner_join(colors, by = c("color_id" = "id"))
```

```
# A tibble: 22 x 5  
  color_id total_batmobile total_batwing name      rgb  
    <dbl>         <dbl>         <dbl> <chr>    <chr>  
1         0           543           418 Black    #05131D  
2         1           33           45 Blue     #0055BF  
3         4           16           81 Red      #C91A09  
4        14           20           22 Yellow   #F2CD37  
5        15           16           22 White    #FFFFFF  
6        36           15            9 Trans-Red #C91A09  
7        57            8            3 Trans-Neon Orange #FF800D  
8        71          202          158 Light Bluish Gray #A0A5A9  
9        72          160          213 Dark Bluish Gray #6C6E68  
10       182            8           14 Trans-Orange #F08F1C  
# ... with 12 more rows
```

Adding percentages

```
batmobile_colors %>%  
  full_join(batwing_colors, by = "color_id", suffix = c("_batmobile", "_batwing")) %>%  
  replace_na(list(total_batmobile = 0, total_batwing = 0)) %>%  
  inner_join(colors, by = c("color_id" = "id")) %>%  
  mutate(total_batmobile = total_batmobile / sum(total_batmobile),  
         total_batwing = total_batwing / sum(total_batwing))
```

```
# A tibble: 22 x 5  
  color_id total_batmobile total_batwing name      rgb  
    <dbl>         <dbl>         <dbl> <chr>    <chr>  
1         0         0.516         0.397 Black    #05131D  
2         1         0.0314        0.0428 Blue     #0055BF  
3         4         0.0152        0.0770 Red      #C91A09  
4        14         0.0190        0.0209 Yellow   #F2CD37  
5        15         0.0152        0.0209 White    #FFFFFF  
6        36         0.0143        0.00856 Trans-Red #C91A09  
7        57         0.00760        0.00285 Trans-Neon Orange #FF800D  
8        71         0.192         0.150  Light Bluish Gray #A0A5A9  
9        72         0.152         0.202  Dark Bluish Gray #6C6E68  
10       182         0.00760        0.0133 Trans-Orange #F08F1C  
# ... with 12 more rows
```

The difference between fractions

```
colors_joined <- batmobile_colors %>%
  full_join(batwing_colors, by = "color_id", suffix = c("_batmobile", "_batwing")) %>%
  replace_na(list(total_batmobile = 0, total_batwing = 0)) %>%
  inner_join(colors, by = c("color_id" = "id")) %>%
  mutate(total_batmobile = total_batmobile / sum(total_batmobile),
         total_batwing = total_batwing / sum(total_batwing),
         difference = total_batmobile - total_batwing)
```

colors_joined

```
# A tibble: 22 x 6
  color_id total_batmobile total_batwing name      rgb      difference
  <dbl>      <dbl>      <dbl> <chr>      <chr>      <dbl>
1         0         0.516         0.397 Black      #05131D      0.119
2         1         0.0314        0.0428 Blue       #0055BF     -0.0114
3         4         0.0152        0.0770 Red        #C91A09     -0.0618
4        14         0.0190        0.0209 Yellow     #F2CD37     -0.00190
5        15         0.0152        0.0209 White      #FFFFFF     -0.00570
6        36         0.0143        0.00856 Trans-Red   #C91A09      0.00570
7        57         0.00760        0.00285 Trans-Neon Orange #FF800D      0.00475
8        71         0.192         0.150 Light Bluish Gray #A0A5A9      0.0418
9        72         0.152         0.202 Dark Bluish Gray #6C6E68     -0.0504
10       182         0.00760        0.0133 Trans-Orange #F08F1C     -0.00570
# ... with 12 more rows
```


Visualizing the data

```
library(ggplot2)
library(forcats)

color_palette <- setNames(colors_joined$rgb, colors_joined$name)

colors_joined %>%
  mutate(name = fct_reorder(name, difference)) %>%
  ggplot(aes(name, difference, fill = name)) +
  geom_col() +
  coord_flip() +
  scale_fill_manual(values = color_palette, guide = FALSE)
```

Visualizing the data



Comparing Batman and Star Wars themes



Let's practice!

JOINING DATA WITH DPLYR